

OTR Bake-off — who wrote the best radio drama


720-word judged rung · 2 BLIND reads (Verne + Codex) · 4 verdicts (Verne, GPT, Antigravity, Grounded-in-repo) · model axis still blocked

◆ CONSENSUS WINNER — EVERY JUDGE, #1

scifi_fable2 — “The Far Shore Relay”

A mirror-bounced wildfire warning forcing a family-vs-valley choice. Best dialogue in the set. **But the win is length-dependent:** at 320/420 it's one long speech like every short-form lane — it only becomes a real multi-character play at 720 (3 → 7 → 64 voiced lines). Not a static edge.

BLIND LEADERBOARD — VERNE + CODEX READ THE SCRIPTS (LOWER IS BETTER)



KEEP / IMPROVE / LEAVE — FOUR VERDICTS

Verne (blind read) + three meta-analyses: GPT, Antigravity, and a **repo-grounded** one (checked against the metrics JSON + the project's own observation/spec docs). **K**=keep **I**=improve **L**=leave. The lone earlier split — science_news — now lands IMPROVE (3 of 4).

BANK	VERNE	GPT	ANTIGRAV	GROUND	NET
scifi_fable2	K	K	K	K	KEEP
public_domain_story	K	K	K	K	KEEP
scifi_codex	K	K	K	K	KEEP
media_archive	K	I	I	K	KEEP
shakespeare	I	I	I	I	IMPROVE
original_radio	I	I	I	I	IMPROVE
scifi_sonnet	I	I	I	I	IMPROVE
science_news	I	L	I	I	IMPROVE

BANK	VERNE	GPT	ANTIGRAV	GROUND	NET
original_codex56sol	I	L	L	L	LEAVE
scifi_gemini	L	L	L	L	LEAVE

NAMED ROOT-CAUSE LEADS — WHERE TO LOOK FIRST (GROUNDED IN THE REPO)

scifi_gemini
leave?

— check FIRST: the spec doc carries **two contradictory engine versions** (archived v2 vs corrected v3). Verify the implementation matches the right one before writing the fragment off as a creative failure — it may be a wiring bug, not the bank.

shakespeare
improve

— the exchange_system "invented modern subtext" fix (2026-07-14) may **not have fully taken** on this leg; its sibling adaptation lane took the same fix clean. Worth a second pass / verify the new-seam leg was judged.

scifi_sonnet
improve

— its spec review flagged **nine tight seams** = heavy surface for the technical model; likely why it's the only bank with an outright 320 FAIL. Seam consolidation is the identified fix.

original_codex56sol
leave

— the 8K-context truncation suspect was **ruled out** (fixed pre-sweep), so the thinness is a genuine content problem (premise doesn't generate story), not a technical one.

scifi_codex
keep

— spec review flagged the exact-word P3 beat table as a re-roll risk and ordered it advisory; confirm that's live.

BEST BANK PER WORD TIER (720 KNOWN · 420 CONFOUNDED · 320 INFERRED)

320W

public_domain

tightest target-hit (1.009) · sonnet FAILED here · fable2 = pure monologue

420W Δ

public_domain

rung CONFOUNDED (3 banks re-legged on looser caps) · fable2 still speech-mode


720W

scifi_fable2

the blind-judged rung — known

Model axis (still blocked): running the mistral/gemma matrices would **isolate prompt vs model** — e.g. is fable2's length-triggered turn into drama aion-specific or general? Are science_news / gemini prompt problems any model hits, or aion-specific? Needs the matrices to answer.

“DIFFERENT SHOWS AT LENGTH” — VOICED LINES ACROSS 320 / 420 / 720



Every lane holds a flat ~14-line shape as the target grows — except **scifi_fable2** (3 → 7 → 64). That single line is the whole story of why it wins only at 720.

MODEL still unanswered: all episodes are aion (creative) + Mistral-Nemo (technical), constant across banks — so every gap above is a BANK effect, not a model one. "Best model" needs the blocked matrices.

Provenance: Verne + Codex = independent BLIND script reads. GPT / Antigravity / Grounded = meta-analyses of the results (rankings mirror Verne's; value is the verdicts + leads). Full grounded writeup: docs/2026-07-15-720-bakeoff-verdict.md.